

In [1]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

In [2]:

```
df = pd.read_csv('logistic_regression.txt')
```


In [3]:

```
df.head()
```

Out[3]:

	loan_amnt	term	int_rate	installment	grade	sub_grade	emp_title	emp_length	home_
0	10000.0	36 months	11.44	329.48	B	B4	Marketing	10+ years	
1	8000.0	36 months	11.99	265.68	B	B5	Credit analyst	4 years	N
2	15600.0	36 months	10.49	506.97	B	B3	Statistician	< 1 year	
3	7200.0	36 months	6.49	220.65	A	A2	Client Advocate	6 years	
4	24375.0	60 months	17.27	609.33	C	C5	Destiny Management Inc.	9 years	N

5 rows × 27 columns



In [4]:

```
df.shape
```

Out[4]:

```
(396030, 27)
```

Their are total 396030 data points, 26 features and 1 label.

In [5]:

```
df.columns
```

Out[5]:

```
Index(['loan_amnt', 'term', 'int_rate', 'installment', 'grade', 'sub_grade',
      'emp_title', 'emp_length', 'home_ownership', 'annual_inc',
      'verification_status', 'issue_d', 'loan_status', 'purpose', 'title',
      'dti', 'earliest_cr_line', 'open_acc', 'pub_rec', 'revol_bal',
      'revol_util', 'total_acc', 'initial_list_status', 'application_type',
      'mort_acc', 'pub_rec_bankruptcies', 'address'],
      dtype='object')
```

displays the columns present in the dataframe

In [6]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 396030 entries, 0 to 396029
Data columns (total 27 columns):
#   Column                Non-Null Count  Dtype
---  -
0   loan_amnt             396030 non-null float64
1   term                  396030 non-null object
2   int_rate              396030 non-null float64
3   installment           396030 non-null float64
4   grade                 396030 non-null object
5   sub_grade            396030 non-null object
6   emp_title             373103 non-null object
7   emp_length           377729 non-null object
8   home_ownership        396030 non-null object
9   annual_inc           396030 non-null float64
10  verification_status    396030 non-null object
11  issue_d               396030 non-null object
12  loan_status           396030 non-null object
13  purpose               396030 non-null object
14  title                 394275 non-null object
15  dti                   396030 non-null float64
16  earliest_cr_line      396030 non-null object
17  open_acc              396030 non-null float64
18  pub_rec               396030 non-null float64
19  revol_bal            396030 non-null float64
20  revol_util           395754 non-null float64
21  total_acc            396030 non-null float64
22  initial_list_status    396030 non-null object
23  application_type       396030 non-null object
24  mort_acc              358235 non-null float64
25  pub_rec_bankruptcies  395495 non-null float64
26  address               396030 non-null object
dtypes: float64(12), object(15)
memory usage: 81.6+ MB
```

We can observe that some features are having object data type which need to be converted to int or float data type.

In [7]:

```
df.describe()
```

Out[7]:

	loan_amnt	int_rate	installment	annual_inc	dti	open_ac
count	396030.000000	396030.000000	396030.000000	3.960300e+05	396030.000000	396030.000000
mean	14113.888089	13.639400	431.849698	7.420318e+04	17.379514	11.311111
std	8357.441341	4.472157	250.727790	6.163762e+04	18.019092	5.137641
min	500.000000	5.320000	16.080000	0.000000e+00	0.000000	0.000000
25%	8000.000000	10.490000	250.330000	4.500000e+04	11.280000	8.000000
50%	12000.000000	13.330000	375.430000	6.400000e+04	16.910000	10.000000
75%	20000.000000	16.490000	567.300000	9.000000e+04	22.980000	14.000000
max	40000.000000	30.990000	1533.810000	8.706582e+06	9999.000000	90.000000

we can observe that the gap between mean and median values is drastic so loan_amnt feature is effected by outliers

int_rate is not effected by outliers as mean and the median values are almost same.

installment is

In [8]:

```
df.describe(include='object')
```

Out[8]:

	term	grade	sub_grade	emp_title	emp_length	home_ownership	verification_status
count	396030	396030	396030	373103	377729	396030	396030
unique	2	7	35	173105	11	6	3
top	36 months	B	B3	Teacher	10+ years	MORTGAGE	Verified
freq	302005	116018	26655	4389	126041	198348	139563

In [9]:

```
df.isnull().sum()
```

Out[9]:

```
loan_amnt      0
term           0
int_rate       0
installment    0
grade          0
sub_grade      0
emp_title      22927
emp_length     18301
home_ownership  0
annual_inc     0
verification_status  0
issue_d        0
loan_status    0
purpose        0
title          1755
dti            0
earliest_cr_line  0
open_acc       0
pub_rec        0
revol_bal      0
revol_util     276
total_acc      0
initial_list_status  0
application_type  0
mort_acc       37795
pub_rec_bankruptcies  535
address        0
dtype: int64
```

we can see emp_title, emp_length, title, revol_util, mort_acc, pub_rec_bankruptcies are having missing values need to impute data using some imputation techniques like mean, median or model based Imputing(KNNImputer).

In [10]:

```
df['loan_status'].value_counts(normalize=True)
```

Out[10]:

```
Fully Paid      0.803871
Charged Off     0.196129
Name: loan_status, dtype: float64
```

the 80.38% people in the data set repayed loan and 19.61% of the data points either repaying the loan or defaulters.

In [11]:

```
df['loan_status'].value_counts()
```

Out[11]:

```
Fully Paid      318357
Charged Off     77673
Name: loan_status, dtype: int64
```

In [12]:

```
df['pub_rec'].value_counts()
```

Out[12]:

```
0.0    338272
1.0     49739
2.0      5476
3.0      1521
4.0       527
5.0       237
6.0       122
7.0        56
8.0        34
9.0         12
10.0        11
11.0         8
13.0         4
12.0         4
19.0         2
86.0         1
40.0         1
17.0         1
15.0         1
24.0         1
Name: pub_rec, dtype: int64
```

most of the people are having good public record only a few are having bad public record.

In [13]:

```
df['mort_acc'].value_counts()
```

Out[13]:

0.0	139777
1.0	60416
2.0	49948
3.0	38049
4.0	27887
5.0	18194
6.0	11069
7.0	6052
8.0	3121
9.0	1656
10.0	865
11.0	479
12.0	264
13.0	146
14.0	107
15.0	61
16.0	37
17.0	22
18.0	18
19.0	15
20.0	13
24.0	10
22.0	7
21.0	4
25.0	4
27.0	3
23.0	2
32.0	2
26.0	2
31.0	2
30.0	1
28.0	1
34.0	1

Name: mort_acc, dtype: int64

only a few number of people are taking loan account in multiple numbers.

In [14]:

```
df['pub_rec_bankruptcies'].value_counts()
```

Out[14]:

```
0.0    350380
1.0     42790
2.0      1847
3.0       351
4.0        82
5.0        32
6.0         7
7.0         4
8.0         2
Name: pub_rec_bankruptcies, dtype: int64
```

In [15]:

```
df['pub_rec'] = df['pub_rec'].apply(lambda x:1 if x>1 else 0)
df['mort_acc'] = df['mort_acc'].apply(lambda x:1 if x>1 else 0)
df['pub_rec_bankruptcies'] = df['pub_rec_bankruptcies'].apply(lambda x:1 if x>1 else 0)
```

converting all the records with value more than 1 as 1 and else 0

In [16]:

```
df['pub_rec'].value_counts()
```

Out[16]:

```
0    388011
1      8019
Name: pub_rec, dtype: int64
```

In [17]:

```
df['mort_acc'].value_counts()
```

Out[17]:

```
0    237988
1    158042
Name: mort_acc, dtype: int64
```

In [18]:

```
df['pub_rec_bankruptcies'].value_counts()
```

Out[18]:

```
0    393705
1     2325
Name: pub_rec_bankruptcies, dtype: int64
```

verifying the changed records are reflected

In [19]:

```
df['term'].value_counts()
```

Out[19]:

```
36 months    302005
60 months     94025
Name: term, dtype: int64
```

In [20]:

```
#df['grade'] = df['grade'].replace(df['grade'].value_counts(normalize=True).index, df['grade']
```

In [21]:

```
df['grade'].replace(df['grade'].value_counts(normalize=True).index, df['grade'].value_counts
```

Out[21]:

```
0      0.292953
1      0.292953
2      0.292953
3      0.162076
4      0.267624
...
396025  0.292953
396026  0.267624
396027  0.292953
396028  0.267624
396029  0.267624
Name: grade, Length: 396030, dtype: float64
```


In [22]:

```
df['sub_grade'].value_counts()
```

Out[22]:

B3	26655
B4	25601
C1	23662
C2	22580
B2	22495
B5	22085
C3	21221
C4	20280
B1	19182
A5	18526
C5	18244
D1	15993
A4	15789
D2	13951
D3	12223
D4	11657
A3	10576
A1	9729
D5	9700
A2	9567
E1	7917
E2	7431
E3	6207
E4	5361
E5	4572
F1	3536
F2	2766
F3	2286
F4	1787
F5	1397
G1	1058
G2	754
G3	552
G4	374
G5	316

Name: sub_grade, dtype: int64

In [23]:

```
#df['emp_title'].replace(df['emp_title'].value_counts(normalize=True).index,df['emp_title'])
```

In [24]:

```
df['emp_title'].value_counts(normalize=True)
```

Out[24]:

```
Teacher          0.011764
Manager          0.011391
Registered Nurse  0.004974
RN               0.004948
Supervisor       0.004905
...
Teachers aide/bus monitor  0.000003
TJCross Engineers  0.000003
assitsant manager  0.000003
Applied Energy    0.000003
Healthcare Call Center Rep  0.000003
Name: emp_title, Length: 173105, dtype: float64
```

In [25]:

```
df['emp_length'] = df['emp_length'].replace(df['emp_length'].value_counts(normalize=True).i
```

converting the feature emp_length to category

In [26]:

```
df['home_ownership'] = df['home_ownership'].replace(df['home_ownership'].value_counts(norma
```

In [27]:

```
df['verification_status'] = df['verification_status'].replace(df['verification_status'].val
```

In [28]:

```
df['issue_d'] = df['issue_d'].replace(df['issue_d'].value_counts(normalize=True).index,df[' '
```

In [29]:

```
#df['loan_status'] = df['loan_status'].replace(df['loan_status'].value_counts(normalize=Tru
```

In [30]:

```
df['purpose'] = df['purpose'].replace(df['purpose'].value_counts(normalize=True).index,df[' '
```

In [31]:

```
#df['title'].replace(df['title'].value_counts(normalize=True).index,df['title'].value_count
```

In [32]:

```
df['title'].value_counts()
```

Out[32]:

```
Debt consolidation          152472
Credit card refinancing    51487
Home improvement           15264
Other                      12930
Debt Consolidation          11608
...
Debt Consolotation         1
My Debt Consolidation loan  1
Short term until 12/31      1
On a Debt Free Adventure of my own!  1
cc debt consolidation       1
Name: title, Length: 48817, dtype: int64
```

In [33]:

```
df['earliest_cr_line'] = df['earliest_cr_line'].replace(df['earliest_cr_line'].value_counts
```

In [34]:

```
df['initial_list_status'] = df['initial_list_status'].replace(df['initial_list_status'].val
```

In [35]:

```
df['application_type'] = df['application_type'].replace(df['application_type'].value_counts
```

In [36]:

```
replacer = dict({'term':{' 36 months':3, ' 60 months':5}, 'grade':{'A':1, 'B':2, 'C':3, 'D':4, 'E
```

In [37]:

```
df['grade']
```

Out[37]:

```
0      B
1      B
2      B
3      A
4      C
..
396025 B
396026 C
396027 B
396028 C
396029 C
Name: grade, Length: 396030, dtype: object
```

In [38]:

```
df = df.replace(replacer)
```

as grade is cardinal and loan status is target variable and term is binary we are converting this using the replace function.

In [39]:

```
df.head()
```

Out[39]:

	loan_amnt	term	int_rate	installment	grade	sub_grade	emp_title	emp_length	home_o
0	10000.0	3	11.44	329.48	2	B4	Marketing	0.333681	
1	8000.0	3	11.99	265.68	2	B5	Credit analyst	0.063411	
2	15600.0	3	10.49	506.97	2	B3	Statistician	0.083989	
3	7200.0	3	6.49	220.65	1	A2	Client Advocate	0.055174	
4	24375.0	5	17.27	609.33	3	C5	Destiny Management Inc.	0.040542	

5 rows × 27 columns

In [40]:

```
df['grade'].value_counts()
```

Out[40]:

```
2    116018
3    105987
1     64187
4     63524
5     31488
6     11772
7      3054
Name: grade, dtype: int64
```

In [41]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 396030 entries, 0 to 396029
Data columns (total 27 columns):
#   Column                Non-Null Count  Dtype
---  -
0   loan_amnt             396030 non-null  float64
1   term                  396030 non-null  int64
2   int_rate              396030 non-null  float64
3   installment           396030 non-null  float64
4   grade                 396030 non-null  int64
5   sub_grade             396030 non-null  object
6   emp_title             373103 non-null  object
7   emp_length           377729 non-null  float64
8   home_ownership        396030 non-null  float64
9   annual_inc            396030 non-null  float64
10  verification_status   396030 non-null  float64
11  issue_d               396030 non-null  float64
12  loan_status           396030 non-null  int64
13  purpose               396030 non-null  float64
14  title                 394275 non-null  object
15  dti                   396030 non-null  float64
16  earliest_cr_line      396030 non-null  float64
17  open_acc              396030 non-null  float64
18  pub_rec               396030 non-null  int64
19  revol_bal             396030 non-null  float64
20  revol_util            395754 non-null  float64
21  total_acc             396030 non-null  float64
22  initial_list_status   396030 non-null  float64
23  application_type      396030 non-null  float64
24  mort_acc              396030 non-null  int64
25  pub_rec_bankruptcies  396030 non-null  int64
26  address               396030 non-null  object
dtypes: float64(17), int64(6), object(4)
memory usage: 81.6+ MB
```

In [42]:

```
df = df.drop(columns=['sub_grade', 'emp_title', 'title', 'address'])
```

In [43]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 396030 entries, 0 to 396029
Data columns (total 23 columns):
#   Column                Non-Null Count  Dtype
---  -
0   loan_amnt              396030 non-null float64
1   term                  396030 non-null int64
2   int_rate              396030 non-null float64
3   installment           396030 non-null float64
4   grade                 396030 non-null int64
5   emp_length            377729 non-null float64
6   home_ownership        396030 non-null float64
7   annual_inc            396030 non-null float64
8   verification_status   396030 non-null float64
9   issue_d               396030 non-null float64
10  loan_status           396030 non-null int64
11  purpose               396030 non-null float64
12  dti                   396030 non-null float64
13  earliest_cr_line      396030 non-null float64
14  open_acc              396030 non-null float64
15  pub_rec               396030 non-null int64
16  revol_bal             396030 non-null float64
17  revol_util            395754 non-null float64
18  total_acc             396030 non-null float64
19  initial_list_status    396030 non-null float64
20  application_type       396030 non-null float64
21  mort_acc              396030 non-null int64
22  pub_rec_bankruptcies  396030 non-null int64
dtypes: float64(17), int64(6)
memory usage: 69.5 MB
```

In [44]:

```
df.isna().sum()
```

Out[44]:

```
loan_amnt      0
term           0
int_rate       0
installment    0
grade          0
emp_length     18301
home_ownership 0
annual_inc     0
verification_status 0
issue_d        0
loan_status    0
purpose        0
dti            0
earliest_cr_line 0
open_acc       0
pub_rec        0
revol_bal      0
revol_util     276
total_acc      0
initial_list_status 0
application_type 0
mort_acc       0
pub_rec_bankruptcies 0
dtype: int64
```

still emp_length and revol_util are having some missing values so we have to fill the missing data with the help of some imputer.

In [45]:

```
df[df['emp_length'].isna()]
```

Out[45]:

	loan_amnt	term	int_rate	installment	grade	emp_length	home_ownership	annual_inc
35	5375.0	3	13.11	181.39	2	NaN	0.403480	34000.00
36	3250.0	3	16.78	115.52	3	NaN	0.403480	22500.00
49	15000.0	3	7.89	469.29	1	NaN	0.500841	90000.00
58	10000.0	3	17.56	359.33	4	NaN	0.500841	32000.00
91	30225.0	5	18.24	771.47	4	NaN	0.500841	65800.00
...
395946	35000.0	5	16.20	854.86	3	NaN	0.500841	84000.00
395963	7000.0	3	20.20	260.86	5	NaN	0.095311	32964.00
395988	35000.0	5	15.59	843.53	4	NaN	0.095311	102396.00
395999	11125.0	3	24.11	437.11	6	NaN	0.500841	31789.80
396015	4000.0	3	9.16	127.50	2	NaN	0.500841	57400.00

18301 rows × 23 columns

In [46]:

```
X = df.drop(columns='loan_status')
Y = df['loan_status']
```

separating the target from the main data

In [47]:

```
from sklearn.impute import SimpleImputer
```

In [48]:

```
imputer = SimpleImputer()
imputer.fit(X)
```

Out[48]:

```
SimpleImputer(add_indicator=False, copy=True, fill_value=None,
               missing_values=nan, strategy='mean', verbose=0)
```

imputing the missing vaalues with the help of mean imputer.

In [49]:

```
X.head()
```

Out[49]:

	loan_amnt	term	int_rate	installment	grade	emp_length	home_ownership	annual_inc
0	10000.0	3	11.44	329.48	2	0.333681	0.403480	117000.0
1	8000.0	3	11.99	265.68	2	0.063411	0.500841	65000.0
2	15600.0	3	10.49	506.97	2	0.083989	0.403480	43057.0
3	7200.0	3	6.49	220.65	1	0.055174	0.403480	54000.0
4	24375.0	5	17.27	609.33	3	0.040542	0.500841	55000.0

5 rows × 22 columns

In [50]:

```
def all_plots(X):  
    for i in X.columns:  
        sns.displot(data=X,x=i)  
        plt.show()
```



function used to plot the univariety plots.



In [51]:



```
plt.rcParams.update({'figure.max_open_warning': 0})
```


In [52]:



```
all_plots(X)
```





























we can see the how the perticular feature is varying some are binary some are multi variet and some are numerical in nature

In [68]:

```
fig, ax = plt.subplots(figsize=(20,20))
sns.heatmap(X.corr(),annot=True,ax=ax)
plt.show()
```



loan_amount and instalment are having high corelation followed by grade and int_rate are having higher corelation pub_rec and pub_rec_bankruptcy ia also having higher corelation

In [49]:

```
X_new = imputer.transform(X)
```

In [50]:

```
from sklearn.preprocessing import StandardScaler
```

In [51]:

```
scaling = StandardScaler()  
scaling.fit(X_new)
```

Out[51]:

```
StandardScaler(copy=True, with_mean=True, with_std=True)
```

standardizing the data using standard scaler

In [52]:

```
from sklearn.model_selection import train_test_split
```

In [53]:

```
x_train,x_test,y_train,y_test = train_test_split(X_new,Y,test_size=0.20,random_state=42)
```

splitting 80% of data to train and remaining 20% data to test

In [54]:

```
x_train.shape,x_test.shape,y_train.shape,y_test.shape
```

Out[54]:

```
((316824, 22), (79206, 22), (316824,), (79206,))
```

In [55]:

```
x_train = scaling.transform(x_train)
```

In [56]:

```
x_test = scaling.transform(x_test)
```

In [57]:

```
from sklearn.linear_model import LogisticRegression
```

importing logistic regression from sklearn

In [58]:

```
model = LogisticRegression(random_state=42).fit(x_train,y_train)
```

In [59]:

```
pred = model.predict(x_test)
```

In [60]:

```
for i,j in zip(X.columns , model.coef_[0]):  
    print(f"The feature {i} feature importance is:- {j}")
```

```
The feature loan_amnt feature importance is:- 0.054416986801077025  
The feature term feature importance is:- -0.21231854222108398  
The feature int_rate feature importance is:- -0.01905039778871478  
The feature installment feature importance is:- -0.1116727229952053  
The feature grade feature importance is:- -0.4384954825225699  
The feature emp_length feature importance is:- 0.02675927582476799  
The feature home_ownership feature importance is:- 0.04253679424952343  
The feature annual_inc feature importance is:- 0.18555571602713264  
The feature verification_status feature importance is:- -0.04099548172404  
5  
The feature issue_d feature importance is:- -0.14512588758930342  
The feature purpose feature importance is:- 0.0074447428705224465  
The feature dti feature importance is:- -0.4360867366819868  
The feature earliest_cr_line feature importance is:- 0.0171016475512284  
The feature open_acc feature importance is:- -0.10054705261849302  
The feature pub_rec feature importance is:- -0.02447002931307284  
The feature revol_bal feature importance is:- 0.06803010744369418  
The feature revol_util feature importance is:- -0.07710480695104927  
The feature total_acc feature importance is:- 0.11389453478009894  
The feature initial_list_status feature importance is:- -0.00389655534819  
70076  
The feature application_type feature importance is:- -0.03057171858511070  
8  
The feature mort_acc feature importance is:- 0.09788161464289102  
The feature pub_rec_bankruptcies feature importance is:- 0.00101710643768  
85601
```

grade is important feature followed by dti

In [61]:

```
from sklearn.metrics import f1_score,precision_score,recall_score,roc_auc_score,confusion_m
```

In [62]:

```
confusion_matrix(y_test.values,pred)
```

Out[62]:

```
array([[ 1226, 14351],  
       [ 1130, 62499]], dtype=int64)
```

In [63]:

```
tn, fp, fn, tp = confusion_matrix(y_test.values,pred).flatten()
```

In [64]:

```
precision_score(y_test.values,pred),tp/(tp+fp)
```

Out[64]:

```
(0.8132595966167859, 0.8132595966167859)
```

our model is giving an precision of 81.32%

In [65]:

```
recall_score(y_test.values,pred),tp/(tp+fn)
```

Out[65]:

```
(0.9822408021499631, 0.9822408021499631)
```

our model is giving an recall of 98.22%

In [66]:

```
f1_score(y_test.values,pred)
```

Out[66]:

```
0.8897984752169363
```

our model is giving an F1 score of 88.97%

In [67]:

```
roc_auc_score(y_test.values,pred)
```

Out[67]:

```
0.5304732931594651
```

In [68]:

```
pred_proba = model.predict_proba(x_test)
```

In [69]:

```
pred_proba
```

Out[69]:

```
array([[0.27767617, 0.72232383],
       [0.39748346, 0.60251654],
       [0.32558046, 0.67441954],
       ...,
       [0.14974418, 0.85025582],
       [0.09429087, 0.90570913],
       [0.40694468, 0.59305532]])
```

predicting the probabilities of both teh class 0 and class 1

In [70]:

```
fpr, tpr, threshold = roc_curve(y_test.values,pred_proba[:,1])
```

In [71]:


```
fpr, tpr, threshold
```

Out[71]:

```
(array([0.00000000e+00, 6.41972138e-05, 6.41972138e-05, ...,
        9.99871606e-01, 1.00000000e+00, 1.00000000e+00]),
 array([0.00000000e+00, 0.00000000e+00, 1.25728834e-04, ...,
        9.99984284e-01, 9.99984284e-01, 1.00000000e+00]),
 array([2.         , 1.         , 0.99795403, ..., 0.24356935, 0.23205378,
        0.18039997]))
```

In [72]:

```
plt.plot(fpr,tpr,marker='.',label='Logistic Regression')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend()
plt.show()
```




plotting the ROC-AUC curve

In [73]:

```
precision_val, recall_val, threshold = precision_recall_curve(y_test,pred_proba[:,1])
```

In [74]:

```
plt.plot(recall_val,precision_val,marker='.',label='Logistic Regression')
plt.xlabel('recall')
plt.ylabel('precision')
plt.legend()
plt.show()
```



plotting the precision recall curve we can observe that if we want to get more precision (less FP) then we have to sacrifice the recall score.

In [75]:

```
max(precision_val),threshold
```

Out[75]:

```
(1.0,
 array([0.18039997, 0.23205378, 0.24101828, ..., 0.99999912, 1.
        1.          ]))
```

In [76]:

```
recall_val
```

Out[76]:

```
array([1.00000000e+00, 9.99984284e-01, 9.99984284e-01, ...,
       1.57161043e-05, 0.00000000e+00, 0.00000000e+00])
```

In [77]:

```
from sklearn.model_selection import GridSearchCV
```

In [78]:

```
model_h_tuned = LogisticRegression(random_state=42)
```

In [96]:

```
hyperparameters = {'C':np.arange(0.1,3.4,0.1)}
best_cv = GridSearchCV(estimator=model_h_tuned,param_grid=hyperparameters,n_jobs=-1,scoring
```


In [97]:

```
best_cv.fit(x_train,y_train)
```

Out[97]:

```
GridSearchCV(cv=None, error_score=nan,
             estimator=LogisticRegression(C=1.0, class_weight=None, dual=False,
             fit_intercept=True,
             intercept_scaling=1, l1_ratio=None
             max_iter=100, multi_class='auto',
             n_jobs=None, penalty='l2',
             random_state=42, solver='lbfgs',
             tol=0.0001, verbose=0,
             warm_start=False),
             iid='deprecated', n_jobs=-1,
             param_grid={'C': array([0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8,
             0.9, 1. , 1.1, 1.2, 1.3,
             1.4, 1.5, 1.6, 1.7, 1.8, 1.9, 2. , 2.1, 2.2, 2.3, 2.4, 2.5, 2.6,
             2.7, 2.8, 2.9, 3. , 3.1, 3.2, 3.3])},
             pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
             scoring='precision', verbose=0)
```

hyperparameter tuning the model to get the best value of C

In [98]:

```
best_cv.best_estimator_
```

Out[98]:

```
LogisticRegression(C=0.8, class_weight=None, dual=False, fit_intercept=True,
                  intercept_scaling=1, l1_ratio=None, max_iter=100,
                  multi_class='auto', n_jobs=None, penalty='l2',
                  random_state=42, solver='lbfgs', tol=0.0001, verbose=0,
                  warm_start=False)
```

In [99]:

```
pred_proba = best_cv.predict_proba(x_test)
```

In [100]:

```
precision_score(y_test.values,best_cv.predict(x_test))
```

Out[100]:

```
0.8132595966167859
```

In [101]:

```
precision_score(y_test.values,pred_proba[:,1]>0.96)
```

Out[101]:

```
0.9625
```

the best possible value of precision is 0.9625 and we are getting this value at the threshold 0.96 instead of 0.5 at the same place our recall score is only 0.84%

In [102]:

```
recall_score(y_test.values, pred_proba[:,1]>0.96)
```

Out[102]:

0.008470980213424696

Actionable Insights & Recommendations

if we want to balance between bad loan and at the same time we have to give loans to the people we are repaying correctly we have to use a threshold which balances out both the precision and recall score which would be somewhere around 0.5 and 0.96.

as grade and dti are contributing the most to the prediction we can fine tune and get the exact values which will definitely help in preventing the bad loans.

more than 80% of the loan given by the bank is repayed so we can improve offer some reduction in the interest for the people who are repaying the loan on time will decrease the bad loan.

In []: